

eda

December 27, 2024

0.1 1. Load and analyse data

Import libraries

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
```

```
[3]: # Loading data
file_path = '../data/raw/transactions_train.csv'
data = pd.read_csv(file_path)

# Check data structure
print(data.info())
print(data.head())

# Ensure date column is in datetime format
data['t_dat'] = pd.to_datetime(data['t_dat'])
print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 31788324 entries, 0 to 31788323
```

```
Data columns (total 5 columns):
```

#	Column	Dtype
0	t_dat	object
1	customer_id	object
2	article_id	int64
3	price	float64
4	sales_channel_id	int64

```
dtypes: float64(1), int64(2), object(2)
```

```
memory usage: 1.2+ GB
```

```
None
```

	t_dat	customer_id	article_id \
0	2018-09-20	000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...	663713001
1	2018-09-20	000058a12d5b43e67d225668fa1f8d618c13dc232df0ca...	541518023
2	2018-09-20	00007d2de826758b65a93dd24ce629ed66842531df6699...	505221004
3	2018-09-20	00007d2de826758b65a93dd24ce629ed66842531df6699...	685687003
4	2018-09-20	00007d2de826758b65a93dd24ce629ed66842531df6699...	685687004

```

      price  sales_channel_id
0  0.050831                2
1  0.030492                2
2  0.015237                2
3  0.016932                2
4  0.016932                2
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31788324 entries, 0 to 31788323
Data columns (total 5 columns):
 #   Column                Dtype
---  -
0   t_dat                 datetime64[ns]
1   customer_id           object
2   article_id            int64
3   price                 float64
4   sales_channel_id      int64
dtypes: datetime64[ns](1), float64(1), int64(2), object(1)
memory usage: 1.2+ GB
None

```

0.2 2. Check data quality

0.2.1 2.1. Check for missing values

```

[4]: # Check for missing values
missing_values = data.isnull().sum()
print(missing_values)

```

```

t_dat          0
customer_id    0
article_id     0
price          0
sales_channel_id  0
dtype: int64

```

The dataset doesn't contain missing values.

0.2.2 2.2. Check for duplicate rows

Check whether we have duplicate rows

```

[5]: # Check for duplicate rows
duplicates = data.duplicated().sum()
print(f'Number of duplicates: {duplicates}')

```

```
Number of duplicates: 2974905
```

As per dataset description, duplicate rows could indicate that a user has bought the same item multiple times.

0.3 2.3 Estimating the distribution of values:

Let's look at the basic statistics for columns:

```
[6]: # Basic statistics on numerical data
print(data.describe())

# Checking for unique values in important columns
print(f"Unique customer_id: {data['customer_id'].nunique()}")
print(f"Unique article_id: {data['article_id'].nunique()}")
```

	t_dat	article_id	price \
count	31788324	3.178832e+07	3.178832e+07
mean	2019-09-15 17:27:46.894452992	6.962272e+08	2.782927e-02
min	2018-09-20 00:00:00	1.087750e+08	1.694915e-05
25%	2019-03-28 00:00:00	6.328030e+08	1.581356e-02
50%	2019-08-25 00:00:00	7.145820e+08	2.540678e-02
75%	2020-03-29 00:00:00	7.865240e+08	3.388136e-02
max	2020-09-22 00:00:00	9.562170e+08	5.915254e-01
std	NaN	1.334480e+08	1.918113e-02

	sales_channel_id
count	3.178832e+07
mean	1.704028e+00
min	1.000000e+00
25%	1.000000e+00
50%	2.000000e+00
75%	2.000000e+00
max	2.000000e+00
std	4.564786e-01
Unique customer_id:	1362281
Unique article_id:	104547

```
[7]: # Step 1: Group by customer_id and date (to define sessions)
# Count the number of items in each session
session_counts = data.groupby(['customer_id', 't_dat']).size().
    ↪reset_index(name='items_per_session')
```

```
[19]: # Step 2: Calculate the distribution of items per session
#item_distribution = session_counts['items_per_session'].value_counts().
    ↪sort_index()
# Calculate the distribution of items per session
# Create a DataFrame that includes count of sessions for each number of items_
    ↪per session
item_distribution = session_counts['items_per_session'].value_counts().
    ↪sort_index().reset_index()
item_distribution.columns = ['items_per_session', 'count_of_sessions']
```

```
print(item_distribution.info())
print(item_distribution.head())
print(item_distribution.tail())
print(item_distribution.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 152 entries, 0 to 151
Data columns (total 2 columns):
#   Column                Non-Null Count  Dtype
---  -
0   items_per_session     152 non-null   int64
1   count_of_sessions     152 non-null   int64
dtypes: int64(2)
memory usage: 2.5 KB
None
```

	items_per_session	count_of_sessions
0	1	2687043
1	2	2120525
2	3	1308634
3	4	883951
4	5	553178

	items_per_session	count_of_sessions
147	199	1
148	220	1
149	221	1
150	336	1
151	570	1

	items_per_session	count_of_sessions
count	152.000000	1.520000e+02
mean	85.480263	5.973802e+04
std	68.304535	3.076907e+05
min	1.000000	1.000000e+00
25%	38.750000	1.750000e+00
50%	76.500000	1.050000e+01
75%	114.250000	5.007500e+02
max	570.000000	2.687043e+06

This DataFrame represents aggregated session data. It shows how many sessions had a specific number of items.

The average number of items per session is approximately 85. The average count of sessions is 59,738

0.4 3. Visualize the distribution of items per session

```
[18]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

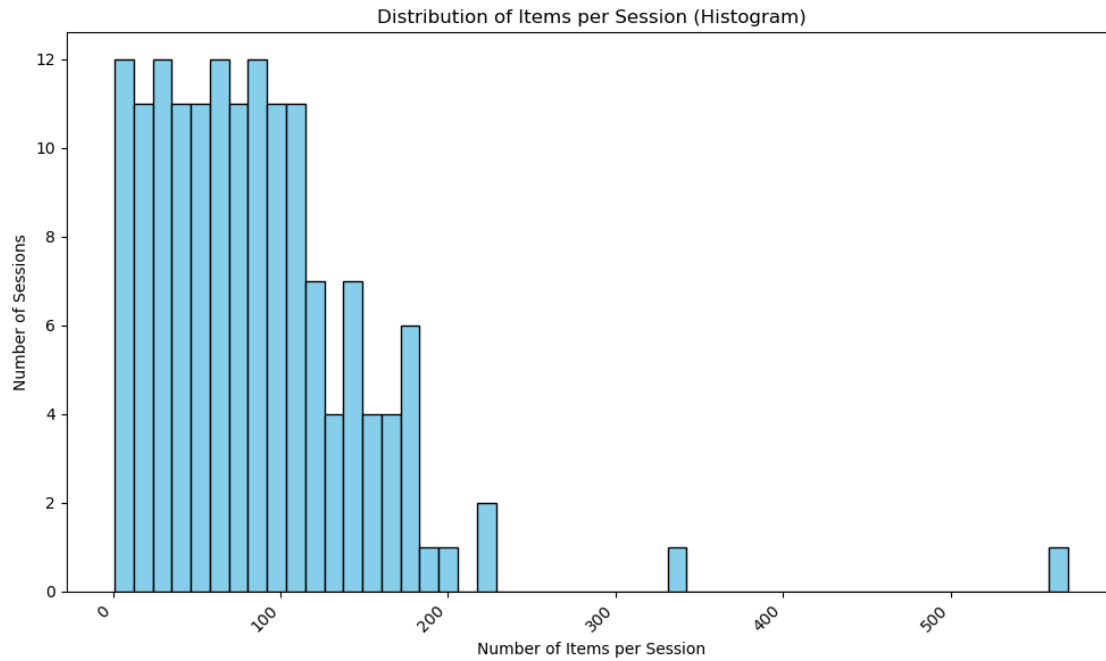
# Assuming your data is in a DataFrame named `item_distribution` and the column
↳with items per session is named `items_per_session`

# Calculate the number of bins
num_bins = 50 # Example: 30 bins

# Create bins
bins = np.linspace(item_distribution['items_per_session'].min(),
↳item_distribution['items_per_session'].max(), num_bins + 1)

# Count items per bin
item_bin_distribution = item_distribution['items_per_session'].
↳value_counts(bins=bins).sort_index()

# Plotting (using histogram for continuous-like data)
plt.figure(figsize=(10, 6))
plt.hist(item_distribution['items_per_session'], bins=bins, edgecolor='black',
↳color='skyblue')
plt.title('Distribution of Items per Session (Histogram)')
plt.xlabel('Number of Items per Session')
plt.ylabel('Number of Sessions')
plt.xticks(rotation=45, ha='right') # Rotate x-axis labels if needed
plt.tight_layout()
plt.show()
```



We see below that most sessions have between 0 and 100 items. The distribution is right-skewed. There is no need to remove outliers as the item count per session will be capped during data preprocessing.